

MSE Database Seminar - Fall 2017

GPU Databases

Author:
Kurath Samuel

Supervisor:
Prof. Stefan Keller

January 22, 2018

Abstract

The content of this paper is splitted into three parts, it begins with an overview of data streams and their difficulties, followed by a part about Apache Storm a distributed stream processing framework. And finally a concrete implementation based on a given problem, solved with Apache Storm. The goal of the implementation is to analyze the minutely updated Augmented Diffs of Open- StreetMap and do a benchmark to test the performance of Apache Storm.

Contents

1. Introduction	2
2. GPU Databases	3
2.1. Initial example	3
2.1.1. General	4
3. MapD	5
3.0.1. Functional Properties	5
3.1. Overview	7
3.1.1. Data Definition Language (DDL)	7
3.1.2. Data Manipulation Language (DML)	8
3.1.3. Data Import and Export	9
3.1.4. Client interfaces	9
3.2. Basics	10
4. Benchmark	11
4.1. NYC Taxi Rides	12
4.2. Queries	12
4.3. Results	12
5. Conclusion	13
Glossary	13
A. Architecture	14
Bibliography	16

1. Introduction

During the Master of Science in Engineering the students have to participate at two seminars. The goal of these is to elaborate a theme on their own, discuss the result in group and write a paper about the topic.

The Databasesystems Seminar does a focus on GPU Database Systems. The students do have a closer look at a certain GPU Database product and have to do a benchmark.

The benchmark is based on the public NYC Taxi Rides dataset and the queries are predetermined by Prof. Stefan Keller.

2. GPU Databases

2.1. Initial example

2. *GPU Databases*

2.1.1. **General**

3. MapD

MapD is a GPU database with the goal to speed up queries and analytic tasks with the power of GPU's and their massive parallel architectures consisting of thousands of cores. The first prototype of MapD was develop in 2012 by Todd Mostak. A year leater MapD was incubated at the MIT's Computer Science and Artificial Intelligence Laboratory (CSAIL) database group and in September 2013 Todd Mostak founded MapD Technologies, Inc. They have got two products called MapD Core [Map17b] and MapD Immerse [Map17c]. The MapD Core SQL engine is an open source in-memory, SQL, GPU database and MapD Immerse is a tool for visual analytics on top of MapD Core SQL engine.



Figure 3.1.: *MapD* Logo

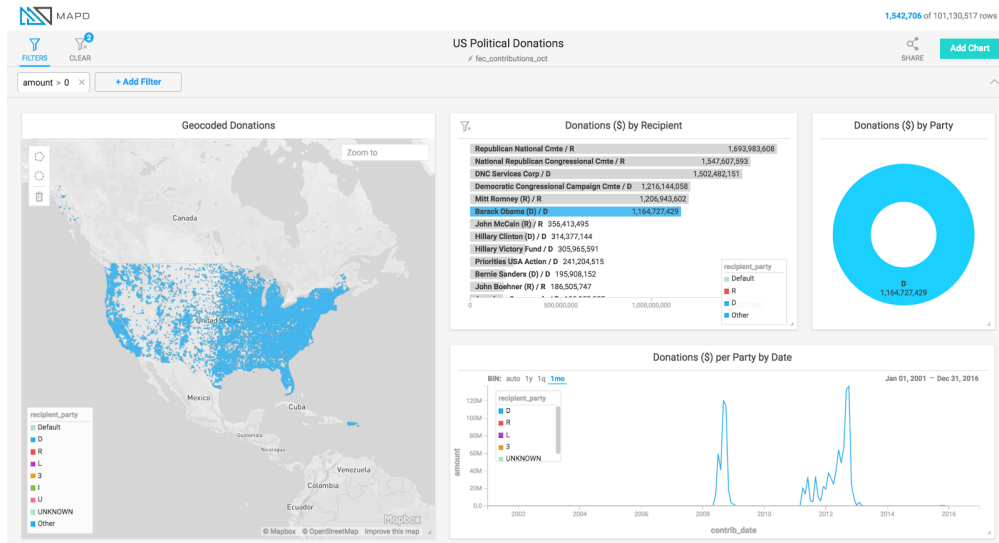


Figure 3.2.: *MapD Immerse* [Dup17]

3.0.1. Functional Properties

Related to the excellent Survey of Sebastian Bress et al. [Bre+14] MapD has the following functional properties.

Storage system MapD has got a relational DBMS that is able to handle data amounts bigger than the memory space. But tries to hold as much data as possible in-memory to to improve the performance.

3. *MapD*

Storage model MapD uses a columnar layout to store the data and uses so called chunks, which split the columns in smaller pieces. The chunks are the basic units of the memory manager.

Processing model MapD processes on operator-at-a-time or one chunk per operation. Thus it is a block-oriented processing. The queries are compiled for the CPU and the GPU.

Query placement In contrast to [Bre+14] the gained experience with MapD showed that MapD tries to run the queries on the GPU even if there isn't enough space and isn't able to handle such queries on his own. °Hence the user had to switch the execution mode from GPU to CPU.

Optimization MapD's optimizer tries to execute the queries on the most suitable device, like text searching using an index on the CPU and table scans on the GPU.

Transactions MapD does not support transactions.

3.1. Overview

The following section will give you an overview about the handling of MapD. Often there will be a comparison between MapD and PostgreSQL to point out the differences and similarities of these two databases.

3.1.1. Data Definition Language (DDL)

The DDL seems familiar since it uses SQL syntax [ddl]. The syntax to handle users, databases, tables, and views is as listed below.

User

- CREATE USER
- DROP USER
- ALTER USER

Database

- CREATE DATABASE
- DROP DATABASE

Table

- CREATE TABLE
- CREATE TABLE AS SELECT
- ALTER TABLE
- DROP TABLE
- TRUNCATE TABLE

View

- CREATE VIEW
- DROP VIEW

3. MapD

Datatypes

MapD supports the data types shown in table 3.1. To get a better intuition the corresponding PostgreSQL data types are listed as well.

MapD		PostgreSQL	
Data type	Size [bytes]	Data type	Size [bytes]
TEXT	Variable	text	Variable
TIMESTAMP	8	timestamp	8
TIME	8	time	8
DATE	8	date	4
FLOAT	4	real	4
DOUBLE	8	double precision	8
INTEGER	4	integer	4
SMALLINT	2	smallint	2
BIGINT	8	bigint	8
BOOLEAN	1	boolean	1
DECIMAL	8	numeric	variable

Table 3.1.: *Data types [Map17a] [Gro17a]*

As you can see MapD supports the common data types. And if you compare them to the corresponding PostgreSQL types they have got nearly the same names. In addition PostgreSQL provides further data types like json or box which allow extended possibilities of use.

3.1.2. Data Manipulation Language (DML)

As well as the DDL of MapD the DML uses the SQL syntax too. MapD currently supports the instructions:

- INSERT INTO
- SELECT

Until now there is no support for the operations:

- DELETE
- UPDATE

But they are in development, since MapD don't want to compromising the speed much with these instructions it may take a while.

Furthermore, MapD provides operations like EXPLAIN, LIKELY/UNLIKELY, Aggregate Functions and Conditional Expressions to improve the DML operations and extend the functionality.

3. MapD

3.1.3. Data Import and Export

MapD allows to import data from different sources as you can see in the following section.

COPY FROM

The COPY FROM operation is callable from the mapdql terminal and imports data from a local CSV or related format file into the database.

SQL Importer

The SQL Importer is a java tool that allows to run queries on other database via JDBC and stores the results in to MapD.

StreamInsert

The StreamInsert program could be attached at the end of a real-time stream processing engine like Kafka or a similar product to import stream data into MapD for further analytic tasks.

HDFS

The tool sqoop-export offers the possibility to import CSV or Parquet files from a HDFS file system into MapD database.

3.1.4. Client interfaces

MapD provides a tool called mapdql [Map17d] as a client-side SQL console that displays query results you submit to the MapD Core Server. The counterpart of PostgreSQL is psql [Gro17b].

Database connection

```
mapdql: mapdql <database> -u <user> -p <password> -port <port> -s <host>
```

```
pgsql: psql -h <host> -p <port> -U <user> -W <password> <database>
```

3. MapD

Commands

The table 3.2 lists some basic commands of mapdql and pgsq. It is only a slight slice of all possible commands, but another good example to point out how much in common those two products have.

Command	mapdql	pgsq
List databases	\l	\l
List tables in database	\t	\d
Describe a table	\d <table>	\d <table>
Connect to a database	\c	\c
Print timing information	\timing	\timing
Switch to GPU mode	\gpu	-
Switch to CPU mode	\cpu	-
Quit	\q	\q

Table 3.2.: *Commands*

3.2. Basics

Installation, mapdql, similar to postgres \t \d etc.

4. Benchmark

Benchmark

4. Benchmark

4.1. NYC Taxi Rides



Figure 4.1.: *Taxi Dropoffs*

4.2. Queries

The following list is about the queries

4.3. Results

The benchmark led to the following results.

5. Conclusion

A. Architecture

Appendices

Bibliography

- [Bre+14] Sebastian Breß et al. “Gpu-accelerated database systems: Survey and open challenges”. In: *Transactions on Large-Scale Data-and Knowledge-Centered Systems XV*. Springer, 2014, pp. 1–35.
- [Dup17] Tai Dupree. *Quick Insight with MapD Immerse Cross Filtering*. Dec. 2017. URL: <https://www.mapd.com/blog/2017/04/14/quick-insight-with-mapd-immersed-cross-filtering/>.
- [Gro17a] The PostgreSQL Global Development Group. *Chapter 8. Data Types*. Dec. 2017. URL: <https://www.postgresql.org/docs/10/static/datatype.html>.
- [Gro17b] The PostgreSQL Global Development Group. *psql*. Dec. 2017. URL: <https://www.postgresql.org/docs/10/static/app-psql.html>.
- [Map17a] Inc. MapD Technologies. *Datatypes and Fixed Encoding*. Dec. 2017. URL: <https://www.mapd.com/docs/latest/mapd-core-guide/fixed-encoding/>.
- [Map17b] Inc. MapD Technologies. *MapD Core - In-Memory, Open Source SQL engine*. Dec. 2017. URL: <https://www.mapd.com/platform/core/>.
- [Map17c] Inc. MapD Technologies. *MapD Immerse Visual Analytics*. Dec. 2017. URL: <https://www.mapd.com/platform/immerse/>.
- [Map17d] Inc. MapD Technologies. *mapdql*. Dec. 2017. URL: <https://www.mapd.com/docs/latest/mapd-core-guide/mapdql/>.